

---

# Amazon Review Sentiment as a Predictor of Economic Trends

## Introduction

This project investigates whether consumer sentiment expressed in Amazon product reviews can serve as an early indicator of broader economic trends. By analyzing millions of reviews across diverse product categories using natural language processing (NLP), we aim to uncover patterns in tone, emotional intensity, and volume—particularly in price-sensitive categories such as groceries, baby products, toys, and media. These patterns are compared to key macroeconomic indicators, including inflation rates and consumer confidence indexes.


We hypothesize that online reviews, authored spontaneously by millions of individuals, may capture real-time fluctuations in public sentiment and financial behavior more quickly than traditional economic reporting. By aggregating sentiment features—such as polarity, subjectivity, and linguistic trends—on an annual basis, we assess whether collective consumer expression can reliably predict periods of economic growth or contraction.

If successful, this approach could offer economists, businesses, and policymakers a high-frequency, low-cost, crowd-sourced signal of economic sentiment—potentially preceding traditional metrics.

---

## Methods


Our modeling process followed a structured pipeline of data preprocessing, feature extraction, aggregation, transformation, and supervised learning using logistic regression. All modeling was conducted with attention to temporal structure, using chronological splits to prevent leakage.



**Figure 1:** Overview of the modeling pipeline, showing major data transformation and modeling stages

We utilized two main datasets: (1) millions of Amazon reviews, filtered to four categories—two “essential” (Groceries and Baby Products) and two “luxury” (Toys & Games and Movies & TV) and (2) a historical GDP per capita dataset for the United States, spanning 1996 to 2023. Reviews with neutral ratings (3 stars) were excluded to sharpen sentiment signals. The remaining reviews included star ratings,

text, category, and timestamps. GDP data was converted into binary labels: 1 for years with growth, 0 otherwise.



**Figure 2:** Trends in key U.S. economic indicators between 1996 and 2023. GDP per capita growth (black) served as the basis for labeling each year as an economic upturn or downturn in our classification task. Additional indicators—such as inflation, unemployment, and interest rates—are shown for context but were not directly used as input features in the model.

## Text Processing and Feature Extraction

We implemented a PySpark NLP pipeline to transform review text into numeric representations. Reviews were tokenized, stop words were removed, and the resulting tokens were mapped into high-dimensional hashed term frequency vectors using the `HashingTF` function with 10,000 buckets. These frequency vectors were then scaled using inverse document frequency (IDF) weighting to emphasize rare but potentially informative terms.

$$c_j^{(d)} = \sum_{t \in d} \mathbf{1}\{h(t) = j\}, \quad j = 1, \dots, D.$$

**Figure 2:** Illustration of text transformation steps from raw reviews to TF-IDF vectors.

In parallel, we performed sentiment scoring using TextBlob. Each review received a polarity score between  $-1$  and  $+1$ , indicating its sentiment intensity. The polarity score was later combined with the TF-IDF representation to form a unified feature vector for each review.

$$\text{idf}(t) = \ln \frac{N_{\text{docs}} + 1}{\text{df}(t) + 1}$$

**Figure 3:** Formula showing how polarity and TF-IDF vectors are concatenated into a final input vector.

### Aggregation by Year and Product Group

To align text and sentiment features with economic indicators, we aggregated data at the year and product category level. For each combination, we computed:

- Mean sentiment polarity
- Sentiment volatility (standard deviation)
- Total review count
- Mean TF-IDF vector

$$\text{tfidf}(t, d) = \text{tf}(t, d) \times \text{idf}(t)$$

**Figure 4:** Yearly aggregation process for sentiment and TF-IDF vectors across essential and luxury categories.

These aggregated features were then pivoted so that each year contained two sets of features—one for essential and one for luxury products—enabling separate modeling of necessity- and discretion-based sentiment.

### Sentiment Statistics and Composite Features

The aggregated sentiment features were further processed to quantify changes in consumer tone and emotional volatility. For each year, the following sentiment features were created:

- `essential_avg_sentiment`, `essential_review_count`, `essential_volatility`
- `luxury_avg_sentiment`, `luxury_review_count`, `luxury_volatility`

These statistics provided macro-style consumer behavior signals, which were then merged with GDP labels and TF-IDF features.

$$\text{polarity}(d) \in [-1, 1]$$

**Figure 5:** Range of polarity values returned by TextBlob for review `d`

We also calculated the sentiment difference between essential and luxury categories, allowing the model to learn contrasts in consumer confidence across economic necessity.

## Dimensionality Reduction

High-dimensional TF-IDF vectors were reduced via Principal Component Analysis (PCA), with 50 components retained for each product category. This yielded 100 features per year, preserving variance while reducing sparsity and complexity.

$$\text{label}(d) = \begin{cases} 1.0, & \text{rating} \geq 4 \\ 0.0, & \text{rating} \leq 2 \\ (\text{discard}), & \text{rating} = 3 \end{cases}$$

**Figure 6:** Label assignment based on star ratings. Neutral reviews (3 stars) were discarded to emphasize clear sentiment signals.

## Year-Over-Year Feature Engineering

Recognizing that raw sentiment levels may be confounded by product trends or seasonal effects, we computed delta-based features that captured year-to-year changes. Specifically, we calculated:

- `essential_sentiment_yoy_change`
- `luxury_sentiment_yoy_change`
- `sentiment_gap_yoy_change`

These features quantify shifts in sentiment and emotional divergence between essential and luxury goods, which may better reflect economic momentum.

$$\mathbf{x}_d = [\text{tfidf\_vec}_d, \text{polarity}(d)].$$

**Figure 7:** Final review-level input vector composed of TF-IDF text features and a polarity score.

Any rows missing valid YOY differences were removed, reducing the usable dataset to 27 years (1997–2023), with 23 years used for training and 5 years for testing.

## Model Setup and Feature Vector

We created the final input feature vector by combining the three year-over-year sentiment features with the 100 PCA components. These features were standardized using `StandardScaler` to prevent uneven feature weighting due to scale differences.

The binary label for each year indicated whether GDP per capita increased (1) or decreased/remained constant (0) compared to the prior year.

$$\mu_y^{(g)} = \frac{1}{|D_{y,g}|} \sum_{d \in D_{y,g}} \text{features}_d$$

**Figure 8:** Aggregation of review features for year  $y$  and group  $g$  (essential or luxury), averaged across all reviews in that group.

## Model Training and Evaluation

We trained a logistic regression model with L1 regularization (`elasticNetParam` = 1.0, `regParam` = 0.01) using a chronological split: 1997–2018 for training, 2019–2023 for testing. This time-respecting strategy avoided leakage and preserved forecasting realism. Performance was measured using AUC and accuracy.

$$\text{sentiment\_gap}_y = \bar{s}_y^{(\text{essential})} - \bar{s}_y^{(\text{luxury})}$$

**Figure 9:** Difference in average sentiment polarity between essential and luxury categories for year  $y$ .

---

## Results

### Data Exploration

Preliminary analysis of the Amazon review data revealed a highly polarized rating distribution. Most reviews clustered at 1, 4, or 5 stars, with a minimal presence of 3-star ratings—validating their exclusion from modeling. Review volume steadily increased over time, particularly after 2012, consistent with Amazon’s platform growth. Sentiment analysis showed that luxury product categories (e.g., Movies & TV) exhibited more fluctuation in polarity than essential categories (e.g., Groceries), suggesting heightened sensitivity to economic conditions in discretionary spending sectors.

Notably, the year 2020 saw a pronounced drop in average sentiment polarity and a spike in polarity volatility across all categories, aligning with the COVID-19 pandemic. These disruptions hinted at a broader link between consumer sentiment and real-world economic stressors.

## Preprocessing and Feature Engineering

TF-IDF vectors derived from review text were hashed into 262,144-dimensional representations and subsequently reduced to 50 principal components per product category using PCA. This dimensionality reduction maintained interpretability while reducing computational load and mitigating overfitting risks. In parallel, sentiment polarity scores were extracted using TextBlob and averaged across each year.

To reflect directional sentiment shifts rather than static scores, year-over-year deltas were computed for each category's average sentiment and their inter-category sentiment gap. This aggregation produced a compact, interpretable dataset containing three engineered sentiment features and 100 PCA components per year.

## Model Performance

The logistic regression model was trained on 23 years of data (1997–2018) and tested on the 5 most recent years (2019–2023). The model achieved:

- **Training Accuracy:** 1.000, **Training AUC:** 1.000
- **Test Accuracy:** 0.800, **Test AUC:** 0.833

$$\text{AUC} = \int_0^1 \text{TPR}(t) \, d[\text{FPR}(t)]$$

**Figure 11.** Area under the ROC curve (AUC), calculated as the integral of the true positive rate (TPR) with respect to the false positive rate (FPR).

$$\text{Accuracy} = \frac{\text{TP} + \text{TN}}{\text{Total samples}}$$

**Figure 12.** Accuracy, defined as the ratio of correct predictions (true positives and true negatives) to the total number of test samples.

The model correctly predicted key economic shifts, including the 2020 downturn, demonstrating its ability to align with actual GDP movement using only review-derived sentiment features. However, not all predictions were correct, and a 0.167 gap between training and test AUC signals some degree of overfitting.

---

## Discussion

The primary objective of this study was to assess whether machine learning models trained solely on Amazon review sentiment could predict U.S. economic upturns and downturns. Specifically, we explored whether aggregated consumer sentiment—captured through polarity scores and textual features—could serve as a timely, crowd-sourced indicator of macroeconomic conditions.

Our dataset consisted of millions of reviews, from which we selected four product categories grouped into two economic segments: “essential” (Groceries and Baby Products) and “luxury” (Toys & Games and Movies & TV). This division was designed to probe whether discretionary spending patterns (luxury goods) responded differently to economic changes compared to necessity-driven purchases (essentials). Exploratory analysis supported this approach: luxury categories exhibited sharper fluctuations in rating distribution and sentiment polarity over time, while essential goods remained more stable.

We retained four key fields from each review: the star rating, review text, timestamp, and product category. Reviews with neutral 3-star ratings were removed to focus the analysis on clearly positive or negative sentiment. TextBlob was used to assign polarity scores to each review, and Spark’s machine learning pipeline was employed to tokenize, filter, and vectorize the text using **HashingTF** and **IDF**. Each review was ultimately represented as a combined sentiment-text vector and assigned a binary label based on its star rating.

To make the data suitable for time-series modeling, we aggregated features annually and by product group. For each year, we computed the average sentiment polarity, polarity volatility (standard deviation), review count, and mean TF-IDF vector for both luxury and essential reviews. Principal component analysis (PCA) reduced the TF-IDF vectors from over 260,000 dimensions to 50 components per category, allowing us to retain the most meaningful patterns while keeping the feature space computationally manageable.

We then calculated annual sentiment deltas—capturing year-over-year changes in average polarity for each category, as well as the change in the sentiment gap between them. These engineered features allowed our model to detect directional trends and momentum rather than static sentiment levels. This approach also aligned the feature timeline with GDP’s annual reporting structure, creating a consistent and interpretable economic time series.

We labeled each year from 1997 to 2023 as either an upturn (GDP per capita increased) or downturn (GDP per capita decreased) using official economic data. The model was trained on data from 1997 to 2018 and tested on data from 2019 to 2023, using a strict chronological split to preserve temporal structure. Randomized cross-validation was intentionally avoided after preliminary experiments showed it yielded poor performance (AUC ~0.58) due to temporal leakage—mixing past and future data disrupts time-dependent patterns, weakening the model’s ability to generalize in real-world forecasting scenarios.

The final logistic regression model achieved perfect training accuracy and AUC (1.000) but still demonstrated strong generalization with 0.800 accuracy and 0.833 AUC on the test set. While the perfect training performance raises overfitting concerns, several factors suggest it may not be overly problematic. First, the model's simplicity—using only three engineered sentiment features and PCA-compressed components—limits its capacity to memorize training data. Second, the strong test performance confirms the model's ability to capture underlying patterns. Third, GDP trends may be linearly separable given the nature of the aggregated sentiment data, especially when annual differences are used.

Nonetheless, the 0.167 drop in AUC from train to test is nontrivial. To address this, future work should consider incorporating additional regularization, exploring ensemble methods, or expanding the feature set to include other metadata (e.g., verified purchase flags, helpful vote counts, review length). Additionally, using non-linear classifiers like Random Forests or XGBoost may help capture interactions that logistic regression misses.

There are also notable limitations. Although our dataset contains millions of reviews, the number of training samples is small—only 27 years of data are available. Amazon was founded in 1995, and reviews only became sufficiently populated after 2012, limiting the statistical depth of our modeling. Furthermore, we analyzed only four product categories. A more granular sectoral breakdown—such as electronics, apparel, or household goods—might uncover more targeted signals tied to economic activity.

Beyond modeling improvements, another key direction is increasing temporal granularity. By aggregating reviews on a quarterly or monthly basis, we could produce more frequent predictions and potentially improve alignment with short-term economic changes. This would also yield more training examples, increasing statistical robustness. Integrating sentiment-based predictors with traditional economic indicators such as unemployment, inflation, or consumer confidence indices may also enhance forecasting accuracy through ensemble approaches.

Finally, one long-term goal is to enable real-time economic tracking using live review data. If consumer sentiment trends can be captured continuously, models like ours could serve as low-latency complements to government-released economic statistics—offering early warnings for economic shifts based on spontaneous consumer expression.

In summary, our results demonstrate that even with a small annual sample and simple sentiment features, Amazon review data contains meaningful signals about macroeconomic activity. With further refinement, this approach holds promise for building high-frequency, crowd-sourced indicators of national economic health.

---

## Conclusion

This study demonstrates that combining sentiment metrics with TF-IDF-based text features from Amazon reviews can provide predictive signals for distinguishing between years of U.S. economic expansion and recession. By aggregating consumer sentiment across both essential and luxury product categories, and



focusing on year-over-year changes rather than absolute values, we were able to extract meaningful macro-level trends from millions of micro-level interactions.

The final logistic regression model achieved perfect training performance (AUC = 1.000, Accuracy = 1.000), but somewhat lower performance on the test set (AUC = 0.833, Accuracy = 0.800). This gap indicates potential overfitting, suggesting that while the model successfully captured patterns in the training data, its generalization to future, unseen years is limited.

Nevertheless, the strong test results affirm that sentiment data—when properly engineered and temporally structured—can offer valuable insight into broader economic patterns. The model’s success in forecasting the 2020 economic downturn, which aligns with COVID-19 disruptions, provides a particularly compelling proof of concept.

To enhance robustness and predictive reliability, future work should explore more aggressive regularization (e.g., increasing `regParam` or tuning `elasticNetParam`), expanding the feature set, or combining sentiment indicators with traditional economic variables in an ensemble framework. Additionally, while our study was constrained to annual granularity, transitioning to quarterly or even monthly analysis could improve temporal resolution and lead to more timely economic forecasts.

Overall, while this sentiment-driven model is not yet a substitute for formal economic forecasting methods, it opens the door to lightweight, scalable tools that tap into real-time consumer behavior for early signals of economic shifts.

---

## Collaboration

- **Ann Nguyen** –
- **Artien Voskanian** –
- **Joanna Tam** –
- **Matthew Mitchell** – GitHub management (hosting the repo, uploading files, and merging branches), README writing (original Module 2 and 3 submissions as well as the final submission), and Final Report Writing (namely Introduction, Method, and Results sections)